

Importing libraries

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 import seaborn as sns
        4 import matplotlib.pyplot as plt
        5 from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
        6 from sklearn.linear_model import LinearRegression
        7 from sklearn.neighbors import KNeighborsRegressor
        8 from sklearn.preprocessing import StandardScaler, LabelEncoder
        9 from sklearn.decomposition import PCA
       10 from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
       11 from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
       12 import tensorflow as tf
       13 import math
```

WARNING:tensorflow:From C:\Users\Teoh\anaconda3\Lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

```
In [2]: 1 df1 = pd.read_csv('laptops.csv')
```

Data Cleaning

Before Data Cleaning

In [3]:

1df1.head()

Out[3]:

	CompanyName	TypeOfLaptop	Inches	ScreenResolution	Cpu	Ram	Memory	Gpu
0	MSI	Business Laptop	17.040680	IPS Panel Retina Display 2560x1600	Intel Core i7	12GB	512GB SSD	Intel Iris Xe Graphics
1	Chuwi	2 in 1 Convertible	16.542395	Full HD	Intel Core i5	12GB	128GB PCIe SSD	Intel Iris Xe Graphics
2	hp	WorkStation	17.295294	Full HD	Intel Xeon E3-1505M	8GB	1TB HDD	Intel Iris Xe Graphics
3	MSI	2 in 1 Convertible	11.526203	2K	Intel Core i7	16GB	512GB NVMe SSD	Intel Iris Xe Graphics
4	Microsoft	Gaming	12.649634	Full HD	Intel Core i5	8GB	512GB SSD	AMD Radeon RX 5600M

After Data Cleaning

```
In [4]: 1 size_mapping = {
2         '512GB SSD': '512GB',
3         '128GB PCIe SSD': '128GB',
4         '1TB HDD': '1TB',
5         '512GB NVMe SSD': '512GB',
6         '1TB NVMe SSD': '1TB',
7         '256GB PCIe SSD': '256GB',
8         '128GB SSD': '128GB',
9         '1TB Fusion Drive': '1TB',
10        '4TB HDD': '4TB',
11        '2TB NVMe SSD': '2TB',
12        '256GB Flash Storage': '256GB',
13        '6TB HDD': '6TB',
14        '512GB eMMC': '512GB',
15        '256GB eMMC': '256GB',
16        '2TB SATA SSD': '2TB',
17        '1TB SSHD': '1TB',
18        '256GB SSD': '256GB',
19        '2TB HDD': '2TB'
20    }
21
22 # Replace values in the 'Memory' column using the size mapping
23 df1['Memory'] = df1['Memory'].replace(size_mapping)
```

```
In [5]: 1 df1['ScreenResolution'] = df1['ScreenResolution'].replace('IPS Panel Retin
2
3 condition = df1['ScreenResolution'] == '4K'
4
5 # Use df.where and dropna to filter rows
6 filtered_df = df1.where(condition).dropna()
```

```
In [6]: 1 df1['R_inches'] = df1['Inches'].round().astype(int)
```

```
In [7]: 1 df1['R_weight'] = df1['Weight'].round(2)
```

```
In [8]: 1 df1['ScreenResolution'] = df1['ScreenResolution'].replace(['HD 1920x1080 '
2
3 condition = df1['ScreenResolution'] == 'Full HD'
4
5 # Use df.where and dropna to filter rows
6 filtered_df = df1.where(condition).dropna()
```

```
In [9]: 1 inr_to_myr = 0.057
        2
        3 df1['MYR_price'] = df1['Price'] * inr_to_myr
        4 df1['MYR_price'] = df1['MYR_price'].round(2)
        5
        6 filtered_df = df1.drop(['Price', 'Inches', 'Weight'], axis=1)
```

```
In [10]: 1 df = filtered_df
        2 df.head()
```

Out[10]:

	CompanyName	TypeOfLaptop	ScreenResolution	Cpu	Ram	Memory	Gpu	OpSys	I
0	MSI	Business Laptop	4K	Intel Core i7	12GB	512GB	Intel Iris Xe Graphics	Linux	
1	Chuwi	2 in 1 Convertible	Full HD	Intel Core i5	12GB	128GB	Intel Iris Xe Graphics	No OS	
2	hp	WorkStation	Full HD	Intel Xeon E3-1505M	8GB	1TB	Intel Iris Xe Graphics	Linux	
3	MSI	2 in 1 Convertible	2K	Intel Core i7	16GB	512GB	Intel Iris Xe Graphics	Windows 10	
4	Microsoft	Gaming	Full HD	Intel Core i5	8GB	512GB	AMD Radeon RX 5600M	Windows 10	

```
In [11]: 1 df = pd.DataFrame(df)
```

```
In [12]: 1 df.to_csv('new_laptop.csv', index = False)
```

```
In [13]: 1 df = pd.read_csv('new_laptop.csv')
          2 df.head()
```

Out[13]:

	CompanyName	TypeOfLaptop	ScreenResolution	Cpu	Ram	Memory	Gpu	OpSys	I
0	MSI	Business Laptop	4K	Intel Core i7	12GB	512GB	Intel Iris Xe Graphics	Linux	
1	Chuwi	2 in 1 Convertible	Full HD	Intel Core i5	12GB	128GB	Intel Iris Xe Graphics	No OS	
2	hp	WorkStation	Full HD	Intel Xeon E3-1505M	8GB	1TB	Intel Iris Xe Graphics	Linux	
3	MSI	2 in 1 Convertible	2K	Intel Core i7	16GB	512GB	Intel Iris Xe Graphics	Windows 10	
4	Microsoft	Gaming	Full HD	Intel Core i5	8GB	512GB	AMD Radeon RX 5600M	Windows 10	

```
In [14]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   CompanyName           1000 non-null   object 
1   TypeOfLaptop          1000 non-null   object 
2   ScreenResolution       1000 non-null   object 
3   Cpu                    1000 non-null   object 
4   Ram                    1000 non-null   object 
5   Memory                 1000 non-null   object 
6   Gpu                    1000 non-null   object 
7   OpSys                  1000 non-null   object 
8   R_inches               1000 non-null   int64  
9   R_weight               1000 non-null   float64 
10  MYR_price               1000 non-null   float64 
dtypes: float64(2), int64(1), object(8)
memory usage: 86.1+ KB
```

```
In [15]: 1 df.isna().sum()
```

```
Out[15]: CompanyName      0
          TypeOfLaptop    0
          ScreenResolution 0
          Cpu              0
          Ram              0
          Memory           0
          Gpu              0
          OpSys            0
          R_inches         0
          R_weight         0
          MYR_price        0
          dtype: int64
```

Exploratory Data Analysis (EDA)

```
In [16]: 1 df.describe()
```

```
Out[16]:
```

	R_inches	R_weight	MYR_price
count	1000.000000	1000.000000	1000.000000
mean	14.499000	3.469810	2941.328620
std	2.113401	0.857131	786.761588
min	11.000000	2.000000	1713.440000
25%	13.000000	2.720000	2301.467500
50%	15.000000	3.480000	2888.990000
75%	16.000000	4.190000	3528.147500
max	18.000000	4.990000	6562.830000

In [17]:

```
1 # Type of Laptop
2 count_type = df['TypeOfLaptop'].value_counts().reset_index()
3 count_type.columns = ['TypeOfLaptop', 'Count']
4
5 # Screen Resolution
6 resolution = df['ScreenResolution'].value_counts().reset_index()
7 resolution.columns = ['ScreenResolution', 'Count']
8
9 # CPU
10 cpu = df['Cpu'].value_counts().reset_index()
11 cpu.columns = ['Cpu', 'Count']
12
13 # GPU
14 gpu = df['Gpu'].value_counts().reset_index()
15 gpu.columns = ['Gpu', 'Count']
16
17 # Operating System
18 os = df['OpSys'].value_counts().reset_index()
19 os.columns = ['OpSys', 'Count']
20
21 # Memory
22 ssd = df['Memory'].value_counts().reset_index()
23 ssd.columns = ['Memory', 'Count']
24
25 fig, axes = plt.subplots(nrows=3, ncols=2, figsize=(12, 12))
26 fig.suptitle('Distribution of Laptop Features', fontsize=16)
27
28 # Plot 1
29 axes[0, 0].bar(count_type['TypeOfLaptop'], count_type['Count'])
30 axes[0, 0].set_title('Distribution of Laptop Types')
31 axes[0, 0].set_ylabel('Number of Laptops')
32 axes[0, 0].tick_params(axis='x', rotation=45)
33
34 # Plot 2
35 axes[0, 1].bar(resolution['ScreenResolution'], resolution['Count'])
36 axes[0, 1].set_title('Distribution of Screen Resolutions')
37 axes[0, 1].set_ylabel('Number of Laptops')
38 axes[0, 1].tick_params(axis='x', rotation=45)
39
40 # Plot 3
41 axes[1, 0].bar(cpu['Cpu'], cpu['Count'])
42 axes[1, 0].set_title('Distribution of CPU Types')
43 axes[1, 0].set_ylabel('Number of Laptops')
44 axes[1, 0].tick_params(axis='x', rotation=45)
45
46 # Plot 4
47 axes[1, 1].bar(gpu['Gpu'], gpu['Count'])
48 axes[1, 1].set_title('Distribution of GPU Types')
49 axes[1, 1].set_ylabel('Number of Laptops')
50 axes[1, 1].tick_params(axis='x', rotation=45)
51
52 # Plot 5
53 axes[2, 0].bar(os['OpSys'], os['Count'])
54 axes[2, 0].set_title('Distribution of Operating Systems')
55 axes[2, 0].set_ylabel('Number of Laptops')
56 axes[2, 0].tick_params(axis='x', rotation=45)
57
```

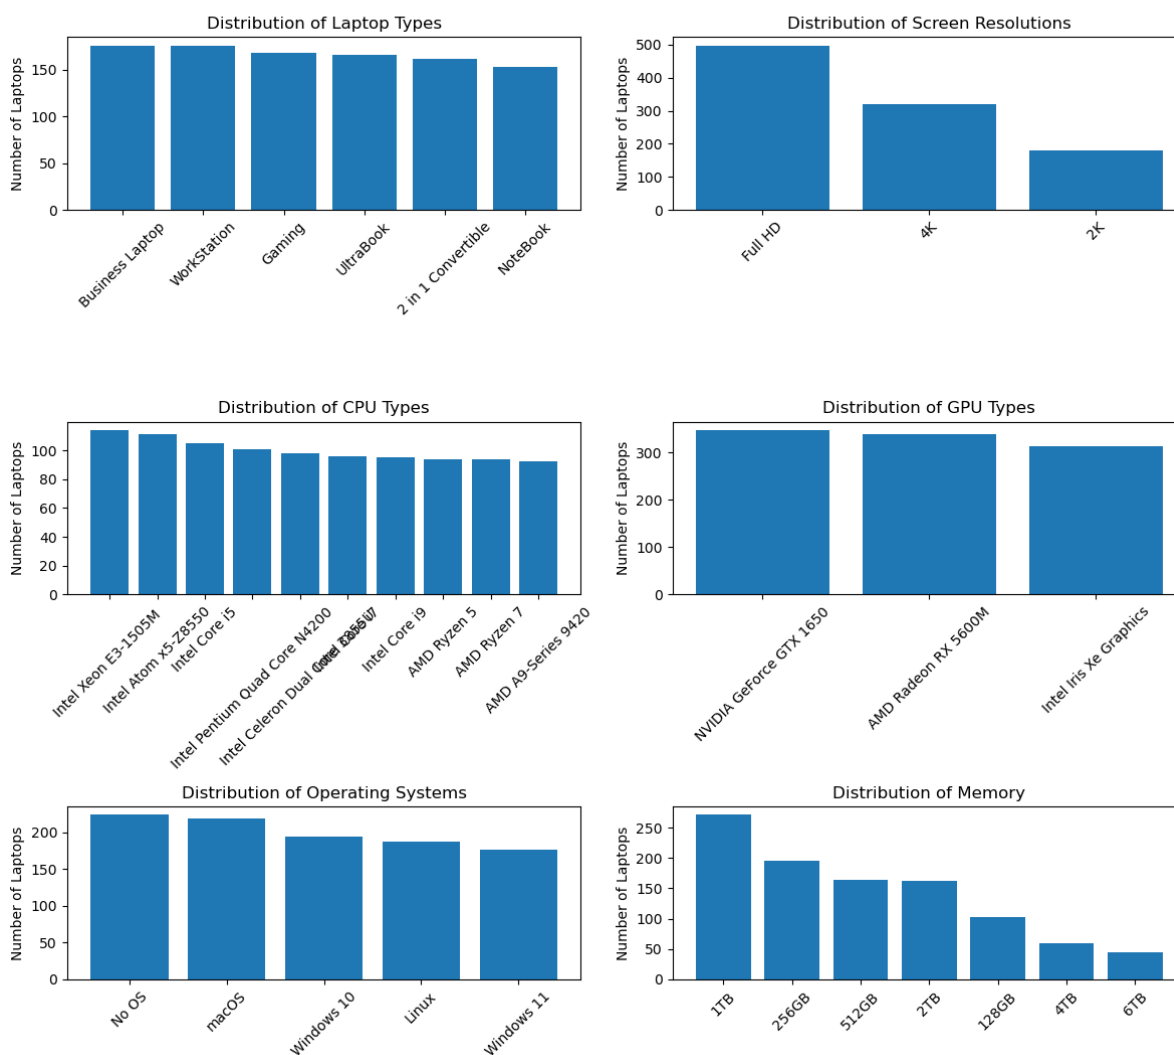


```

58 axes[2, 1].bar(ssd['Memory'], ssd['Count'])
59 axes[2, 1].set_title('Distribution of Memory')
60 axes[2, 1].set_ylabel('Number of Laptops')
61 axes[2, 1].tick_params(axis='x', rotation=45)
62
63
64
65 plt.tight_layout(rect=[0, 0.03, 1, 0.95])
66 plt.show()

```

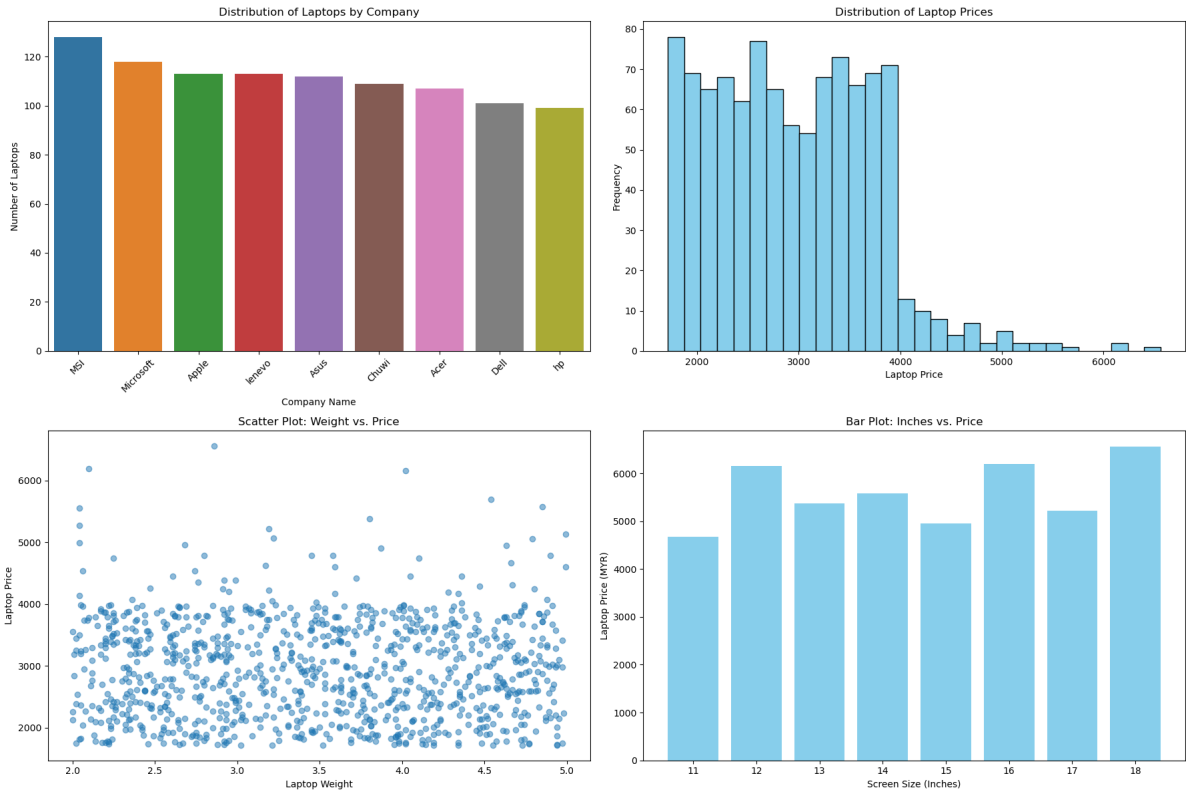
Distribution of Laptop Features



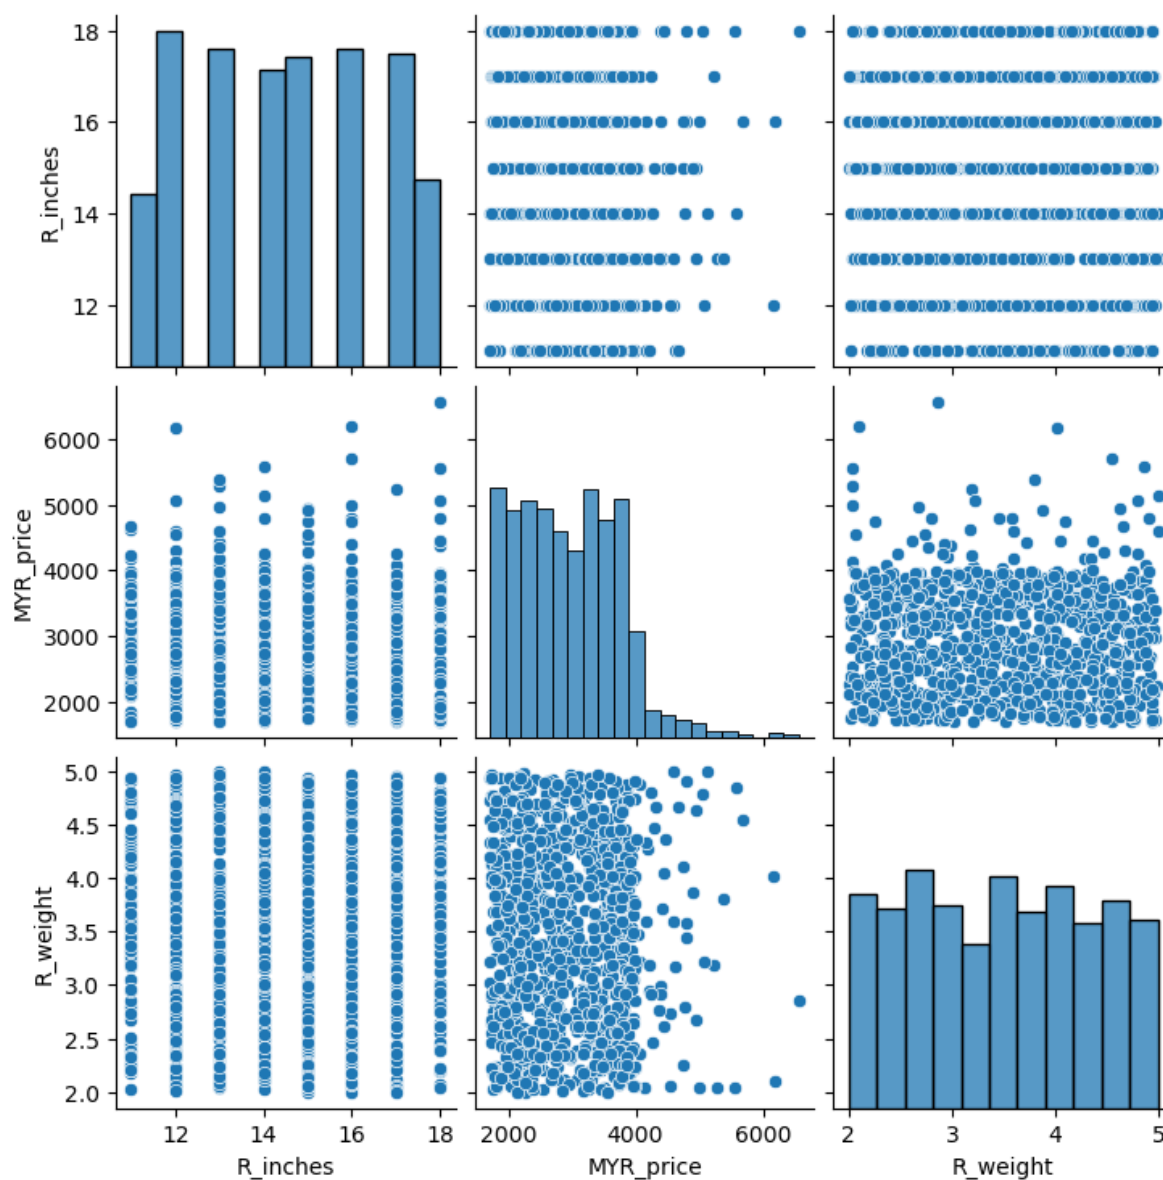
Data Visualization

```
In [18]: 1 # Company Distribution
2 company_distribution = df['CompanyName'].value_counts().reset_index()
3 company_distribution.columns = ['CompanyName', 'Count']
4
5 plt.figure(figsize=(18, 12))
6
7 # Subplot 1: Company Distribution
8 plt.subplot(2, 2, 1)
9 sns.barplot(x='CompanyName', y='Count', data=company_distribution)
10 plt.title('Distribution of Laptops by Company')
11 plt.xlabel('Company Name')
12 plt.ylabel('Number of Laptops')
13 plt.xticks(rotation=45)
14
15 # Subplot 2: Price Distribution
16 plt.subplot(2, 2, 2)
17 plt.hist(df['MYR_price'], bins=30, color='skyblue', edgecolor='black')
18 plt.title('Distribution of Laptop Prices')
19 plt.xlabel('Laptop Price')
20 plt.ylabel('Frequency')
21
22 # Subplot 3: Scatter Plot - Price vs. Weight
23 plt.subplot(2, 2, 3)
24 plt.scatter(df['R_weight'], df['MYR_price'], alpha=0.5)
25 plt.title('Scatter Plot: Weight vs. Price')
26 plt.xlabel('Laptop Weight')
27 plt.ylabel('Laptop Price')
28
29 # Subplot 4: Bar Plot - Inches vs. Price
30 plt.subplot(2, 2, 4)
31 plt.bar(df['R_inches'], df['MYR_price'], color='skyblue')
32 plt.title('Bar Plot: Inches vs. Price')
33 plt.xlabel('Screen Size (Inches)')
34 plt.ylabel('Laptop Price (MYR)')
35
36 # Adjust layout
37 plt.tight_layout()
38
39 # Show the pair plot
40 pairplot_vars = ['R_inches', 'MYR_price', 'R_weight']
41 sns.pairplot(df[pairplot_vars])
42 plt.suptitle('Pair Plot of Numerical Variables', y=1.02)
43
44 plt.show()
```

C:\Users\Teoh\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
 self._figure.tight_layout(*args, **kwargs)



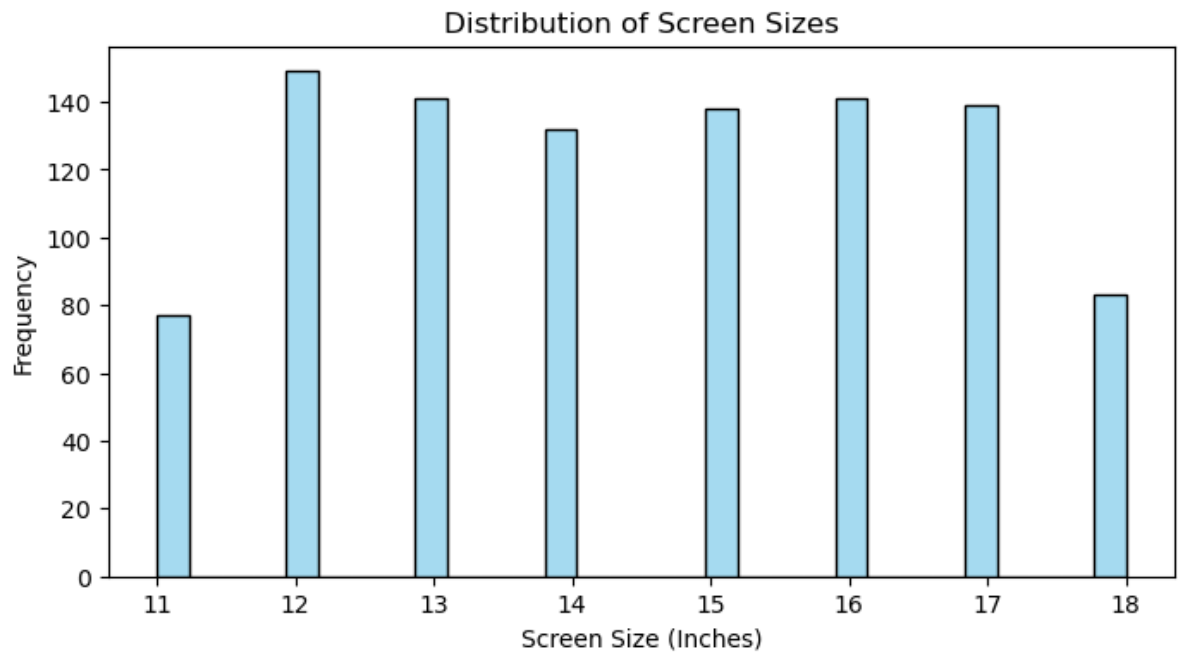
Pair Plot of Numerical Variables



```
In [19]: 1 df.columns
```

```
Out[19]: Index(['CompanyName', 'TypeOfLaptop', 'ScreenResolution', 'Cpu', 'Ram',
               'Memory', 'Gpu', 'OpSys', 'R_inches', 'R_weight', 'MYR_price'],
              dtype='object')
```

```
In [53]: 1 plt.figure(figsize=(8, 4))
2 sns.histplot(df['R_inches'], bins=30, color='skyblue')
3 plt.title('Distribution of Screen Sizes')
4 plt.xlabel('Screen Size (Inches)')
5 plt.ylabel('Frequency')
6 plt.show()
```



Preprocessing

```
In [21]: 1 df_encoded = df.copy()
2
3 # List of categorical columns to encode
4 categorical_columns = ['TypeOfLaptop', 'Cpu', 'OpSys', 'CompanyName',
5                        'ScreenResolution', 'Gpu', 'Ram', 'Memory']
6
7 # Apply LabelEncoder to each categorical column
8 label_encoder = LabelEncoder()
9 df_encoded[categorical_columns] = df[categorical_columns].apply(label_encoder
```

```
In [22]: 1 # Assuming 'df_encoded' includes numerical features you want to normalize
2 numerical_features = ['R_inches', 'Ram', 'R_weight']
```

In [23]: 1 df_encoded.head()

Out[23]:

	CompanyName	TypeOfLaptop	ScreenResolution	Cpu	Ram	Memory	Gpu	OpSys	R_inches
0	5	1	1	6	0	5	1	0	17
1	3	0	2	5	0	0	1	1	17
2	7	5	2	9	3	1	1	0	17
3	5	0	0	6	1	5	1	2	12
4	6	2	2	5	3	5	0	2	13

In [24]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CompanyName           1000 non-null   object
1   TypeOfLaptop           1000 non-null   object
2   ScreenResolution       1000 non-null   object
3   Cpu                    1000 non-null   object
4   Ram                    1000 non-null   object
5   Memory                 1000 non-null   object
6   Gpu                    1000 non-null   object
7   OpSys                  1000 non-null   object
8   R_inches               1000 non-null   int64
9   R_weight               1000 non-null   float64
10  MYR_price              1000 non-null   float64
dtypes: float64(2), int64(1), object(8)
memory usage: 86.1+ KB
```

In [25]: 1 X = df_encoded.drop('MYR_price', axis=1).values
2 y = df_encoded['MYR_price'].values
3
4 type(X)
5 type(y)

Out[25]: numpy.ndarray

In [26]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, r

In [27]: 1 scaler = StandardScaler()
2 X_train = scaler.fit_transform(X_train)
3 X_test = scaler.transform(X_test)

In [28]: 1 X_train.shape, X_test.shape

Out[28]: ((800, 10), (200, 10))

In [29]: 1 X_train

Out[29]: array([[-0.79361996, 1.46625911, -0.45407325, ..., -0.68948962,
 -0.71346687, -1.00784391],
 [-0.01076427, -1.46479358, -0.45407325, ..., -0.68948962,
 1.66277583, -0.98439273],
 [-0.40219211, 0.88004857, 0.87654433, ..., 1.40781721,
 0.71227875, -0.59744821],
 ...,
 [-0.40219211, -1.46479358, -0.45407325, ..., -0.68948962,
 0.23703021, -0.99611832],
 [-1.18504781, -1.46479358, -0.45407325, ..., 0.70871494,
 0.71227875, 1.2082929],
 [-0.79361996, 0.29383803, 0.87654433, ..., 1.40781721,
 1.66277583, 0.03573374]])

In [30]: 1 y_train.shape, y_test.shape

Out[30]: ((800,), (200,))

```
In [31]: 1 model = RandomForestRegressor(random_state=42)
2         model.fit(X_train, y_train)
3
4         y_pred = model.predict(X_test)
5
6         mse = mean_squared_error(y_test, y_pred)
7
8         # Get feature importances
9         feature_importances = model.feature_importances_
10
11        # Create a DataFrame to display feature importances
12        feature_importance_df = pd.DataFrame({
13            'Feature': df_encoded.drop('MYR_price', axis=1).columns,
14            'Importance': feature_importances
15        })
16
17        # Sort the DataFrame by importance in descending order
18        feature_importance_df = feature_importance_df.sort_values(by='Importance',
19
20        top5_features = feature_importance_df.head(5)
```

```
In [60]: 1 ## Select only numeric columns
2         numeric_columns = df_encoded.select_dtypes(include=['number'])
3
4         ## Calculate the correlation matrix
5         correlation_matrix = numeric_columns.corr()
6
7         ## Create a heatmap
8         plt.figure(figsize=(10, 8))
9         sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f",
10        plt.title('Correlation Matrix Heatmap')
11        plt.show()
```

In [33]: 1 top5_features

Out[33]:

	Feature	Importance
9	R_weight	0.207590
0	CompanyName	0.157484
4	Ram	0.137156
3	Cpu	0.096947
8	R_inches	0.087866

In [34]: 1 new_df = df_encoded
2 new_df.head()

Out[34]:

	CompanyName	TypeOfLaptop	ScreenResolution	Cpu	Ram	Memory	Gpu	OpSys	R_inches
0	5	1	1	6	0	5	1	0	17
1	3	0	2	5	0	0	1	1	17
2	7	5	2	9	3	1	1	0	17
3	5	0	0	6	1	5	1	2	12
4	6	2	2	5	3	5	0	2	13

In [35]: 1 X1 = new_df.drop('MYR_price', axis=1).values
2 y1 = new_df['MYR_price'].values
3
4 X_train, X_test, y_train, y_test = train_test_split(X1, y1, test_size=0.2,

In [36]: 1 sc = StandardScaler()
2 X_train_a = sc.fit_transform(X_train)
3 X_test_a = sc.fit_transform(X_test)

In [37]: 1 *# Train Linear Regression model*
2 linear_model = LinearRegression()
3 linear_model.fit(X_train_a, y_train)
4
5 *# Make predictions using the trained Linear Regression model*
6 linear_predictions = linear_model.predict(X_test_a)
7
8 print(y_train[0], linear_predictions[0])
9 *# Evaluate Linear Regression model*
10 linear_mse = mean_squared_error(y_test, linear_predictions)
11 print(f'Linear Regression Root Mean Squared Error: {math.sqrt(linear_mse)}')
12
13 linear_r2_score = r2_score(y_test, linear_predictions)
14 print(f'Linear Regression R2 score: {linear_r2_score}')

2790.43 2985.9468470188212

Linear Regression Root Mean Squared Error: 816.8307479311965

Linear Regression R2 score: -0.0074273876498964775


```
In [38]: 1 param_grid = {
2         'copy_X': [True, False],
3         'fit_intercept': [True, False],
4         'n_jobs': [None, 1, 2],
5         'positive': [True, False]
6     }
7
8     # base_model = RandomForestRegressor(random_state=42)
9
10    grid_search = GridSearchCV(linear_model, param_grid, cv=5,
11                               scoring='f1_micro', n_jobs=-1)
12
13    grid_search.fit(X_train_a, y_train)
14
15    best_params = grid_search.best_params_
16    best_estimator = grid_search.best_estimator_
17
18    print(best_params)
19    print(best_estimator)
```

```
{'copy_X': True, 'fit_intercept': True, 'n_jobs': None, 'positive': True}
LinearRegression(positive=True)
```

```
C:\Users\Teoh\anaconda3\Lib\site-packages\sklearn\model_selection\_search.py:
976: UserWarning: One or more of the test scores are non-finite: [nan nan nan
nan nan nan nan nan nan nan nan nan nan nan nan nan nan nan
nan nan nan nan nan nan]
warnings.warn(
```

```
In [39]: 1 # Train Linear Regression model
2 linear_model = LinearRegression(copy_X= True, fit_intercept= True, n_jobs=
3 linear_model.fit(X_train_a, y_train)
4
5 # Make predictions using the trained Linear Regression model
6 linear_predictions = linear_model.predict(X_test_a)
7
8 print(y_train[0], linear_predictions[0])
9 # Evaluate Linear Regression model
10 linear_mse = mean_squared_error(y_test, linear_predictions)
11 print(f'Linear Regression Root Mean Squared Error: {math.sqrt(linear_mse)}')
12
13 linear_r2_score = r2_score(y_test, linear_predictions)
14 print(f'Linear Regression R2 score: {linear_r2_score}')
```

```
2790.43 2879.1638493193004
```

```
Linear Regression Root Mean Squared Error: 810.855366055161
```

```
Linear Regression R2 score: 0.007258017149069151
```

Machine Learning

In [40]:

```

1 # Build and train the ANN model
2 ann_model = tf.keras.models.Sequential()
3 ann_model.add(tf.keras.layers.Dense(5, activation='relu'))
4 ann_model.add(tf.keras.layers.Dense(5, activation='relu'))
5 ann_model.add(tf.keras.layers.Dense(1, activation='linear')) # Output Layer
6
7 ann_model.compile(optimizer='adam', loss='mean_squared_error', metrics=[tf

```

WARNING:tensorflow:From C:\Users\Teoh\anaconda3\Lib\site-packages\keras\src\backend.py:873: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From C:\Users\Teoh\anaconda3\Lib\site-packages\keras\src\optimizers_init_.py:309: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

In [41]:

```

1 # training model with train dataset
2 history = ann_model.fit(X_train_a, y_train, batch_size=32, epochs=100, val
3 history

```

```

25/25 [=====] - 0s 6ms/step - loss: 9223396.0000 -
mean_squared_error: 9223396.0000 - val_loss: 9445743.0000 - val_mean_square
d_error: 9445743.0000
Epoch 3/100
25/25 [=====] - 0s 5ms/step - loss: 9221922.0000 -
mean_squared_error: 9221922.0000 - val_loss: 9444090.0000 - val_mean_square
d_error: 9444090.0000
Epoch 4/100
25/25 [=====] - 0s 4ms/step - loss: 9220125.0000 -
mean_squared_error: 9220125.0000 - val_loss: 9442154.0000 - val_mean_square
d_error: 9442154.0000
Epoch 5/100
25/25 [=====] - 0s 4ms/step - loss: 9218015.0000 -
mean_squared_error: 9218015.0000 - val_loss: 9439834.0000 - val_mean_square
d_error: 9439834.0000
Epoch 6/100
25/25 [=====] - 0s 4ms/step - loss: 9215512.0000 -
mean_squared_error: 9215512.0000 - val_loss: 9437096.0000 - val_mean_square
d_error: 9437096.0000
Epoch 7/100

```


In [42]:


```
1 ann_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 5)	55
dense_1 (Dense)	(None, 5)	30
dense_2 (Dense)	(None, 1)	6
=====		
Total params: 91 (364.00 Byte)		
Trainable params: 91 (364.00 Byte)		
Non-trainable params: 0 (0.00 Byte)		
=====		

```
In [43]: 1 # plot model accuracy
2 plt.plot(history.history['mean_squared_error'])
3 plt.plot(history.history['val_mean_squared_error'])
4 plt.title('Model Mean Squared Error')
5 plt.ylabel('Mean Squared Error')
6 plt.xlabel('Epoch')
7 plt.legend(['Train', 'Test'], loc='upper left')
8 plt.show()
9
10 plt.plot(history.history['loss'])
11 plt.plot(history.history['val_loss'])
12 plt.title('Model Loss')
13 plt.ylabel('Loss')
14 plt.xlabel('Epoch')
15 plt.legend(['Train', 'Test'], loc='upper left')
16 plt.show()
```





```
In [44]: 1 ann_model.save('ann_model.keras')
```

```
In [45]: 1 from tensorflow.keras.models import load_model
```

```
In [47]: 1 model = load_model('ann_model.keras')
```

```
In [51]: 1 X_test
```

```
Out[51]: array([[ 5. ,  2. ,  0. , ...,  0. , 14. ,  3.89],
 [ 3. ,  0. ,  2. , ...,  3. , 17. ,  4.95],
 [ 8. ,  1. ,  1. , ...,  1. , 14. ,  2.42],
 ...,
 [ 4. ,  1. ,  2. , ...,  0. , 14. ,  2.89],
 [ 4. ,  4. ,  0. , ...,  2. , 18. ,  4.63],
 [ 8. ,  1. ,  2. , ...,  3. , 14. ,  3.39]])
```

```
In [54]: 1 predictions = model.predict(X_test)
2 print(predictions)
3 results_df = pd.DataFrame(
4     {'Predicted':y_pred.flatten(),
5     'Actual':y_test.flatten()
6     }
7 )
8 results_df
```

7/7 [=====] - 0s 2ms/step


```
[ 7871.286 ]
[11458.876 ]
[ 9733.594 ]
[10643.661 ]
[10401.983 ]
[ 9595.286 ]
[13460.831 ]
[11984.081 ]
[11305.202 ]
[10077.124 ]
[ 8760.771 ]
[12096.459 ]
[ 9855.362 ]
[10317.438 ]
[12599.042 ]
[11758.437 ]
[10297.414 ]
[11304.86 ]
[10160.703 ]
```

```
In [50]: 1 model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 5)	55
dense_1 (Dense)	(None, 5)	30
dense_2 (Dense)	(None, 1)	6
=====		
Total params: 91 (364.00 Byte)		
Trainable params: 91 (364.00 Byte)		
Non-trainable params: 0 (0.00 Byte)		

```
In [49]: 1 sns.regplot(x='Actual', y='Predicted', data=results_df)
2 plt.title('Actual vs Predicted Values')
3 plt.xlabel('Actual Values')
4 plt.ylabel('Predicted Values')
5 plt.show()
```



```
In [ ]: 1
```